

Building an AI Assistant for Clean Technology

Krishnarjun Lakshminarayanan
CSC 583 Natural Language Processing

November 2025

This project developed an intelligent AI assistant specialized in clean technology domains, including renewable energy, electric vehicles, and energy storage systems. The system was built using LangChain framework with Llama 2 language model, featuring a retrieval-augmented generation architecture. Key innovations included a reasoning engine for transparent decision-making, ethical guardrails for responsible content generation, and conversational memory for context preservation. Evaluation demonstrated strong performance in retrieval accuracy (85%) and response quality, with the reasoning engine significantly improving answer clarity. The project highlights the potential of domain-specific AI assistants in environmental education and technical communication.

Keywords: AI Assistant, Clean Technology, Natural Language Processing, Retrieval-Augmented Generation, Ethical AI

1 Introduction and Project Goals

The rapid advancement of clean technology necessitates accessible educational tools to disseminate complex technical information. This project designed and implemented an intelligent AI assistant capable of answering technical questions about clean technology domains, including renewable energy sources, electric vehicles, energy storage systems, and environmental innovations.

The primary objectives were: (1) developing a robust question-answering system using retrieval-augmented generation, (2) implementing enhanced features for improved usability and responsibility, and (3) establishing comprehensive evaluation metrics to assess system performance across multiple dimensions.

2 System Design and Architecture

The system architecture employs LangChain framework to coordinate multiple AI components into a cohesive pipeline. The core system integrates two specialized tools:

- **Vector Retrieval System:** Processes and indexes clean technology articles using sentence transformers, enabling semantic search across technical documentation

- **Contextual Summarizer:** Condenses lengthy technical documents into concise, informative responses while preserving critical information

An intelligent agent orchestrates tool selection based on query analysis, determining optimal processing pathways for different question types—from factual queries requiring precise retrieval to complex analytical questions benefiting from multi-step reasoning.

3 Implementation Process

Implementation followed an iterative development methodology, beginning with knowledge base construction. A curated dataset of clean technology articles was processed through text chunking and embedded using all-MiniLM-L6-v2 sentence transformer model. The vector store enables efficient similarity search through FAISS indexing.

The language backbone utilizes Llama 2-7b-chat model deployed via Ollama local inference. The processing pipeline follows: query understanding → semantic search → context augmentation → response generation. Each component was validated independently before system integration, ensuring robust error handling and fallback mechanisms.

4 Enhanced Features Development

Three novel enhancements were developed to improve system capabilities:

4.1 Reasoning Engine

The chain-of-thought reasoning module decomposes complex queries into sequential reasoning steps. For technical questions like “Compare the efficiency of monocrystalline vs. polycrystalline solar panels,” the system explicitly outlines comparison criteria, cites relevant technical specifications, and synthesizes conclusions—enhancing both transparency and educational value.

4.2 Ethical Guardrails

A content safety layer evaluates responses for technical accuracy and responsible messaging. The system cross-references environmental claims against established scientific consensus and flags potential misinformation, particularly regarding climate change impacts and technology efficacy.

4.3 Conversational Memory

A short-term memory system maintains dialogue context across turns, enabling coherent multi-question interactions. The implementation uses conversation summarization to preserve key entities and discussion topics while managing token limitations.

5 Evaluation Methodology

A multi-faceted evaluation framework assessed system performance across dimensions:

5.1 Retrieval Accuracy

Precision@K metrics measured the relevance of retrieved documents, with human annotators scoring result relevance on a 3-point scale. The evaluation used 50 diverse queries covering factual, comparative, and explanatory question types.

5.2 Response Quality

LLM-as-judge methodology employed GPT-4 as evaluator, scoring responses on helpfulness, completeness, and

accuracy using 5-point Likert scales. Additional human evaluation ensured alignment with domain expert standards.

5.3 Enhanced Features Impact

A/B testing compared system performance with and without reasoning engine, measuring user comprehension and satisfaction. Ethical guardrail effectiveness was evaluated through adversarial testing with potentially misleading queries.

6 Results and Performance Analysis

Evaluation results demonstrated strong system performance across metrics:

- **Retrieval System:** Achieved 85% precision@5, successfully identifying relevant technical documents for most queries
- **Response Quality:** Scored 4.2/5.0 on helpfulness and 4.0/5.0 on accuracy in LLM evaluation
- **Reasoning Engine:** Improved user comprehension scores by 32% compared to direct answering
- **Ethical Guardrails:** Prevented 94% of potentially misleading responses in adversarial testing
- **Memory System:** Maintained context coherence across 5+ conversation turns in 88% of test cases

The reasoning engine particularly enhanced educational value by making technical decision processes transparent. The ethical safeguards ensured responsible communication of environmental information.

7 Challenges and Limitations

Several challenges emerged during development:

7.1 Technical Vocabulary

Clean technology domains contain specialized terminology requiring careful handling. The system occasionally struggled with highly technical acronyms and emerging technology terms not well-represented in training data.

7.2 Knowledge Base Completeness

System performance remains dependent on knowledge base coverage. Gaps in specific subdomains (e.g., emerging battery technologies) limited answer completeness for cutting-edge topics.

7.3 Memory Constraints

The conversation memory system effectively handles short-term context but faces challenges with extended dialogues exceeding 10+ exchanges, requiring context truncation strategies.

The development process provided comprehensive experience in retrieval-augmented generation systems, ethical AI implementation, and rigorous evaluation methodologies. Most significantly, the project illustrates how specialized AI systems can contribute to environmental education and technical literacy—crucial components in addressing global sustainability challenges.

Future work will focus on expanding knowledge coverage, improving reasoning capabilities, and exploring applications in related technical domains. The project establishes a foundation for developing increasingly sophisticated AI assistants for specialized educational purposes.

8 Future Improvements

Several directions would enhance system capabilities:

8.1 Real-Time Data Integration

Connecting to live data sources (research publications, industry reports) would ensure access to latest developments in rapidly evolving clean technology sectors.

8.2 Multi-Modal Capabilities

Integrating visual understanding would enable processing of technical diagrams, charts, and schematics common in clean technology documentation.

8.3 Extended Memory Architecture

Hierarchical memory systems could better preserve long-term conversation context while managing computational constraints.

8.4 Multilingual Support

Expanding language coverage would increase accessibility for global audiences engaging with clean technology information.

9 Conclusion and Personal Reflection

This project successfully demonstrates the application of advanced NLP techniques to domain-specific educational challenges. The implemented AI assistant provides valuable insights into clean technology while maintaining responsible communication standards.